XUANLEI ZHAO

xuanlei@comp.nus.edu.sg ◇ +65 81485063 ◇ Singapore
Homepage ◇ GitHub ◇ Google Scholar ◇ Twitter

EDUCATION

National University of Singapore   01.2024 - Present
Ph.D. in Computer Science
Supervisor: Yang You

National University of Singapore   08.2022 - 12.2023
M.S. in Computer Science

Huazhong University of Science and Technology   09.2018 - 06.2022
B.Eng. in Computer Science & Electronic Information

RESEARCH INTEREST

- **Machine Learning System**: Parallelism, Scheduling, Offloading, Compiler.
- **Efficient Video Generation**: Efficient Training and Inference, Algorithm-System Co-Design.

RESEARCH EXPERIENCE

PAB: The **First Real-Time** and **Most Cited** cache-based video generation acceleration method.
- **[ICLR 2025]** Real-Time Video Generation with Pyramid Attention Broadcast
- *Xuanlei Zhao*, Xiaolong Jin*, Kai Wang*†, Yang You†*

VideoSys: The **First** and **Most Starred** open-source project for system speedup of video training and inference.
- VideoSys: An Easy and Efficient System for Video Generation
- Project lead.

DCP: The **First Practical** parallel method for efficient variable sequences training (*e.g.,* videos)
- Training Variable Sequences with Data-Centric Parallel
- *Geng Zhang*, Xuanlei Zhao*, Kai Wang†, Yang You†*

DSP: The **Most Efficient** sequence parallel for multi-dim transformers (*e.g.,* spatial-temporal video models).
- DSP: Dynamic Sequence Parallelism for Multi-Dimensional Transformers
- *Xuanlei Zhao, Shenggan Cheng, Chang Chen, Zangwei Zheng, Ziming Liu, Zheming Yang, Yang You*

HeteGen: Accelerate LLM offloading inference by heterogeneous computing between CPU and GPU.
- **[MLSys 2024]** HeteGen: Heterogeneous Parallel Inference for Large Language Models on Resource-Constrained Devices
- *Xuanlei Zhao*, Bin Jia*, Haotian Zhou*, Ziming Liu, Shenggan Cheng, Yang You*

AutoChunk: A compiler to reduce activation memory by over 80% for long sequences (*e.g.,* videos).
- **[ICLR 2024]** AutoChunk: Automated Activation Chunk for Memory-Efficient Long Sequence Inference
- *Xuanlei Zhao, Shenggan Cheng, Guangyang Lu, Jiarui Fang, Haotian Zhou, Bin Jia, Ziming Liu, Yang You*

FastFold: The **First** and **Most Cited** system optimization method for AlphaFold by parallel and computing.
- **[PPoPP 2024]** FastFold: Optimizing AlphaFold Training and Inference on GPU Clusters
- *Shenggan Cheng, Xuanlei Zhao, Guangyang Lu, Jiarui Fang, Tian Zheng, Ruidong Wu, Xiwen Zhang, Jian Peng, Yang You*

INDUSTRY EXPERIENCE

Pika, Inc.   05.2024 - 08.2024
Research Intern | Supervised by Chenlin Meng   *Palo Alto, CA*
· Optimize distributed system on thousands of GPUs for efficient large-scale training of video models.
· Improve training performance with hybrid parallel, I/O optimization, and dynamic activation checkpointing.
· Improve generation efficiency with sequence parallel, adaptive computing, efficient kernel and distillation.

HPC-AI TECH, Inc. (Colossal-AI)   05.2022 - 12.2023
Research Intern | Supervised by Jiarui Fang   *Singapore*
· Contribute 48k lines of code as a core contributor (rank 5th by 2023) and help it gain 35k stars on Github.
· Propose AutoChunk, a compiler to reduce the activation memory by 80% for long sequences inference.
· Participate in the development of various parallelism strategies including sequence parallel, tensor parallel, ZeRO, offloading, auto parallelism and efficient kernels.