# XUANLEI ZHAO

xuanlei@comp.nus.edu.sg ◇ +65 81485063 ◇ Singapore

[Homepage](#) ◇ [GitHub](#) ◇ [Google Scholar](#) ◇ [Twitter](#)

## EDUCATION

**National University of Singapore**     01.2024 - Present
Ph.D. in Computer Science
Supervisor: Yang You

**National University of Singapore**     08.2022 - 12.2023
M.S. in Computer Science

**Huazhong University of Science and Technology**     09.2018 - 06.2022
B.Eng. in Computer Science & Electronic Information

## RESEARCH INTEREST

- **Machine Learning System**: Parallelism, Scheduling, Offloading, Compiler.
- **Efficient Video Generation**: Efficient Training and Inference, Algorithm-System Co-Design.

## RESEARCH HIGHLIGHT

**PAB**: The **First Real-Time** and **Most Cited** cache-based video generation acceleration method.
- Real-Time Video Generation with Pyramid Attention Broadcast
- *Xuanlei Zhao\*, Xiaolong Jin\*, Kai Wang\*[†], Yang You[†]*

**VideoSys**: The **First** and **Most Starred** open-source project for system speedup of video training and inference.
- VideoSys: An Easy and Efficient System for Video Generation
- Project lead.

**DCP**: The **First Practical** parallel method for efficient variable sequences training (*e.g.,* videos)
- Training Variable Sequences with Data-Centric Parallel
- *Geng Zhang\*, Xuanlei Zhao\*, Kai Wang[†], Yang You[†]*

**DSP**: The **Most Efficient** sequence parallel for multi-dim transformers (*e.g.,* spatial-temporal video models).
- DSP: Dynamic Sequence Parallelism for Multi-Dimensional Transformers
- *Xuanlei Zhao, Shenggan Cheng, Chang Chen, Zangwei Zheng, Ziming Liu, Zheming Yang, Yang You*

**HeteGen**: Accelerate LLM offloading inference by heterogeneous computing between CPU and GPU.
- [MLSys 2024] HeteGen: Heterogeneous Parallel Inference for Large Language Models on Resource-Constrained Devices
- *Xuanlei Zhao\*, Bin Jia\*, Haotian Zhou\*, Ziming Liu, Shenggan Cheng, Yang You*

**AutoChunk**: A compiler to reduce activation memory by over 80% for long sequences (*e.g.,* videos).
- [ICLR 2024] AutoChunk: Automated Activation Chunk for Memory-Efficient Long Sequence Inference
- *Xuanlei Zhao, Shenggan Cheng, Guangyang Lu, Jiarui Fang, Haotian Zhou, Bin Jia, Ziming Liu, Yang You*

**FastFold**: The **First** and **Most Cited** system optimization method for AlphaFold from parallel and computing.
- [PPoPP 2024] FastFold: Optimizing AlphaFold Training and Inference on GPU Clusters
- *Shenggan Cheng, Xuanlei Zhao, Guangyang Lu, Jiarui Fang, Tian Zheng, Ruidong Wu, Xiwen Zhang, Jian Peng, Yang You*

## INDUSTRY EXPERIENCE

**Pika, Inc.**                                                      05.2024 - 08.2024

*Research Intern | Supervised by Chenlin Meng*                        *Palo Alto, CA*

· Optimize distributed system on thousands of GPUs for efficient large-scale training of video models.

· Improve training performance with hybrid parallel, I/O optimization, and dynamic activation checkpointing.

· Improve generation efficiency with sequence parallel, adaptive computing, efficient kernel and distillation.

**HPC-AI TECH, Inc. (Colossal-AI)**                                05.2022 - 12.2023

*Research Intern | Supervised by Jiarui Fang*                         *Singapore*

· Contribute 48k lines of code as a core contributor (rank 5th by 2023) and help it gain 35k stars on Github.

· Propose AutoChunk, a compiler to reduce the activation memory by 80% for long sequences inference.

· Participate in the development of various parallelism strategies including sequence parallel, tensor parallel, ZeRO, offloading, auto parallelism and efficient kernels.